

Jonathan Grady Ong

gradyjonathan55@gmail.com | github.com/jong47 | (949) 338-8755 | Irvine, CA

SUMMARY

Backend and AI platform engineer with experience building event-driven Azure systems, LLM gateways, and high-volume document-processing pipelines. Delivered production systems that reduced inference costs 80%, automated 5,000+ client cases, and supported 10+ internal applications using Python, FastAPI, Azure, Redis, and Databricks.

WORK EXPERIENCE

AI Engineer, Tax Relief Advocates – Irvine, CA

03/2026 – Present

- Built a centralized LiteLLM gateway for 10+ internal applications, replacing direct Azure AI Foundry integrations with private-VNet routing, Redis-backed quota-aware load balancing, warm-model failover, and OpenTelemetry cost tracing.
- Architected and helped develop a real-time lead-alert system to reduce time-to-first outreach, enforcing single-owner delivery with Azure Service Bus and delivering Intune desktop alerts via Microsoft Graph and Entra ID within 2 seconds p95 of assignment.
- Reduced OCR and LLM inference costs 80% for a 100K+ page/week document pipeline by routing standard documents to asynchronous Databricks regex extraction and escalating corrupted or malformed files to Azure Document Intelligence and AI Foundry.
- Trained a supervised regression model on 10,000+ historical client cases to predict client service fees from 16–24 intake signals at first client contact, replacing manual fee estimation, with approximately 80% of predictions landing within $\pm 10\%$ of actual fees on a held-out dataset.

Junior AI Engineer, Tax Relief Advocates – Irvine, CA

03/2025 – 03/2026

- Saved \$80K annually by building an automated client-letter web app (React, TypeScript, FastAPI, Google Document AI, Gemini Platform) that cleared a manual backlog of 5,000+ client cases.
- Decreased executive dashboard SQL query latency by 87% (5 min to 40 sec) by reverse-engineering read-only database views and optimizing a query that returned 200K+ rows from a 1.32M row, 200+ column wide table.
- Reduced average tax record audit time by 17% building a horizontally scalable record-merging and deduplication pipeline with data lineage tracing (Celery, Redis, Python, Vertex AI).

AI Engineer Intern, First Databank – Remote

06/2024 – 08/2024

- Improved retrieval accuracy 37% for LLM-powered case study generator using domain-specific keyword filtering for dense vector construction and serving paragraph-level context from Azure Cosmos DB vector search.

TECHNICAL SKILLS

Languages: Python, SQL, TypeScript, Go

Backend & Distributed Systems: FastAPI, REST APIs, asynchronous processing, event-driven architecture, Celery, Redis, Azure Service Bus

AI/ML Platform: Azure AI Foundry, LiteLLM, Azure Document Intelligence, Vertex AI, LLM gateways, RAG, OpenTelemetry

Cloud & Infrastructure: Azure, Azure Blob Storage, Azure Entra ID, Microsoft Graph, Azure VNet, Databricks, Docker, Kubernetes, CI/CD

Data: PostgreSQL, Redis, SQL optimization, data pipelines, Cosmos DB

CERTIFICATIONS

AWS Certified Solutions Architect – Associate

04/2024

PROJECTS

OpenTelemetry Integration for LiteLLM Usage Tracking, Open Source Contribution

02/2026

Authored a feature PR for qodo-ai/pr-agent adding OpenTelemetry instrumentation for LiteLLM usage tracking, capturing token usage, model routing, and completion metrics for monitoring tools like Datadog and Honeycomb.

Meta Research Collaboration, Georgia Institute of Technology

12/2024

Investigated per-token early exit in Llama LayerSkip models with a research team led by Meta AI researcher. Tracked layer-wise token prediction via top-k entropy, analyzing if mid-layer token switching reflected token instability or candidate removal.

EDUCATION

Georgia Institute of Technology, MS in Computer Science

Expected 12/2028

California State University — Fullerton, BS in Computer Science

12/2023